

1

1 / 1 point

Why has "Responsible AI" become a critical area of focus in recent years?



☒ Because algorithms are increasingly used in high-stakes scenarios that directly influence people's lives.

- ☐ Because AI models have become too complex for human oversight.
- ☐ Because regulatory bodies have mandated strict ethical guidelines for all AI development.
- ☐ Because computational power has only recently allowed for fairness checks.

2

1 / 1 point

In the healthcare example discussed, why did the algorithm for identifying patients for managed care inadvertently disadvantage Black patients?



☒ The algorithm used healthcare spending as a proxy for sickness, and Black patients, despite being equally sick, had historically lower healthcare spending due to access issues.

- ☐ Black patients were inherently less likely to benefit from managed care.
- ☐ The algorithm was explicitly coded to consider race as a factor.

3

1 / 1 point

The example of the video interview scoring application, where wearing glasses or a scarf altered personality scores, was used primarily to illustrate:

- ☐ The necessity of diverse training data for facial recognition software.
- ☐ The advanced capabilities of AI in nuanced human assessment.
- ☐ The importance of high-quality video for accurate AI analysis.
- ☒ The need for users to be skeptical of AI applications, especially if they seem too good to be true.

4

1 / 1 point

What does the phrase "human in the loop" refer to in the context of responsible AI?

- ☐ The process of humans manually labeling all training data for AI models.
- ☐ Allowing users to provide direct feedback to the AI system to improve its performance over time.
- ☒ Combining human judgment and oversight with the quantitative outputs of AI models to make final decisions.
- ☐ The ethical requirement for human programmers to be accountable for AI-driven errors.

5

1 / 1 point

The module notes that "today's decisions become tomorrow's training data." What is the primary concern associated with this feedback loop in the context of algorithmic bias?

- ☐ It reduces the model's ability to generalize to new, unseen populations.
- ☐ It makes models overly sensitive to recent trends.
- ☐ It increases the computational cost of retraining models.
- ☒ It can amplify existing biases, as biased decisions fed back into the system will reinforce and potentially worsen the bias in future iterations.

6

1 / 1 point

When discussing strategies to combat bias, what is the first step recommended before even starting to run models?

- ☐ Implementing robust real-time model monitoring systems.
- ☐ Immediately collecting as much diverse data as possible.
- ☐ Selecting the most advanced and fair algorithm available.
- ☒ Pausing to review the problem scope, considering who will be affected and if the right problem is being solved inclusively.

7

1 / 1 point

If you were considering adopting an AI model from a vendor, which question would be most pertinent to ask regarding its applicability to your specific population?

- ☐ "What programming language was used to develop this model?"
- ☐ "How many parameters does your model have?"
- ☒ "Will the model be retrained to fit my organization's population, and how will you monitor its fairness for my specific groups?"
- ☐ "Can you provide a discount if we commit to a long-term contract?"

8

1 / 1 point

What is the core idea behind "individual fairness" as a concept in AI?

- ☐ Ensuring that the AI model performs with equal accuracy across all demographic groups.
- ☐ Allowing each individual to customize the AI's decision-making process for themselves.
- ☐ Guaranteeing that every individual affected by an AI decision has the right to an explanation.
- ☒ Treating similar individuals similarly, irrespective of their demographic characteristics, when they possess comparable relevant attributes.

9

1 / 1 point

"Group fairness" in AI is primarily concerned with:

- ☐ The subjective feeling of fairness experienced by the largest group affected by the AI.
- ☐ The model provides identical outcomes for all individuals within a specific protected group.
- ☒ How the model's decisions and performance metrics (e.g., error rates like false positives) compare across different demographic groups.
- ☐ Ensuring that the development team for the AI is demographically diverse.

10

1 / 1 point

What was the key finding from the experiment where ChatGPT was asked to complete sentences starting with "why doesn't he..." versus "why doesn't she..."?

- ☐ The completions were nearly identical, showing the LLM's neutrality.
- ☐ The LLM provided more creative and complex completions for the female pronoun.
- ☒ The autocompletions were significantly different, reflecting underlying patterns and potential biases in the training data concerning gender.
- ☐ The LLM refused to complete sentences that could perpetuate stereotypes.